

main

March 28, 2019

1 Data Analytics ECS784P

1.1 London House Price Data Analysis and Prediction

1.1.1 Machine Learning of London house price data based on features in yearly data ranging 1995-2018:

- House Price
- House Sales Volume
- House Price Index
- Bank Rate (England)
- GDP (UK)
- GDP Growth (UK)
- Total number household
- Affordability (House price to residence-based earnings ratio)
- Median Income

We are using LinearRegression and GradientBoostingRegressionTree as our methods with the implementation of ShuffleSplit cross-validation along with GridSearchCV for parameters tuning.

1.1.2 Load the data

JS code to prevent sub/scrollable window inside the notebook

```
In [1]: %%javascript
        IPython.OutputArea.prototype._should_scroll = function(lines) {
            return false;
        }
```

<IPython.core.display.Javascript object>

Import external libraries

```
In [2]: import numpy as np
        import pandas as pd
        import matplotlib.pyplot as plt
        import seaborn as sns
        import mpl_toolkits
```

```

from sklearn import preprocessing, model_selection, metrics, svm, ensemble
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import ShuffleSplit, train_test_split, GridSearchCV, learning_curve
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.naive_bayes import GaussianNB
from sklearn.decomposition import PCA
from IPython.display import HTML
from scipy import sparse
from pprint import pprint

%matplotlib inline

```

HTML and JS code to implement toggle to hide or show warning messages that might come up in notebook

```

In [3]: HTML('''<script>
code_show_err=false;
function code_toggle_err() {
  if (code_show_err){
    $('div.output_stderr').hide();
  } else {
    $('div.output_stderr').show();
  }
  code_show_err = !code_show_err
}
$( document ).ready(code_toggle_err);
</script>
To toggle on/off Warning Message, click <a href="javascript:code_toggle_err()">here</a>

```

Out[3]: <IPython.core.display.HTML object>

```

In [4]: # get the data from csv and set it to variable 'data'
data = pd.read_csv('data_year_preprocessed.csv')
data = data.set_index('year')

```

```

In [5]: # peek the data
data

```

```

Out[5]:

```

	price	salesVolume	indexPrice	rate	gdp	gdpGrowth	\
year							
1995	74721.32688	106850.0	18.548293	6.3750	1336125	0.025	
1996	78166.22435	132953.0	19.403430	5.9375	1410855	0.025	
1997	89943.92010	154343.0	22.327042	7.2500	1553949	0.043	
1998	102110.61160	145942.0	25.347216	6.2500	1641822	0.033	
1999	115686.25990	169076.0	28.717139	5.5000	1668683	0.032	
2000	142006.69330	149269.0	35.250737	6.0000	1651392	0.035	
2001	159225.79230	162744.0	39.525084	4.0000	1626218	0.028	
2002	187395.75390	173993.0	46.517796	4.0000	1775814	0.025	
2003	211100.92160	153784.0	52.402199	3.7500	2045693	0.033	

2004	228614.03650	163797.0	56.749531	4.7500	2404700	0.023
2005	235329.26910	136836.0	58.416472	4.5000	2527843	0.031
2006	251281.17760	171450.0	62.376261	5.0000	2700951	0.025
2007	287114.01550	165571.0	71.271151	5.5000	3085300	0.025
2008	282959.01210	80921.0	70.239743	2.0000	2934747	-0.003
2009	257853.75650	75231.0	64.007792	0.5000	2403357	-0.042
2010	284543.12810	91933.0	70.632972	0.5000	2455309	0.017
2011	290551.08050	90041.0	72.124344	0.5000	2635799	0.016
2012	303927.32910	93988.0	75.444769	0.5000	2677082	0.014
2013	329167.88190	115459.0	81.710306	0.5000	2755356	0.020
2014	386124.33550	126356.0	95.848773	0.5000	3036310	0.029
2015	425134.16170	120358.0	105.532298	0.5000	2897060	0.023
2016	467502.88020	104805.0	116.049616	0.2500	2660687	0.018
2017	480240.15790	96023.0	119.211428	0.5000	2622434	0.018
2018	478444.51250	NaN	118.765689	0.7500	2822817	0.014

	numberHousehold	affordability	medianincome
year			
1995	2869079.452	NaN	15247.90251
1996	2887000.000	NaN	15709.96016
1997	2855900.000	NaN	16114.26061
1998	2872500.000	NaN	16114.26061
1999	2901200.000	NaN	16114.26061
2000	2923900.000	NaN	16600.00000
2001	2963800.000	NaN	17600.00000
2002	2971400.000	6.90	18500.00000
2003	2979500.000	7.44	18800.00000
2004	2963300.000	7.95	19100.00000
2005	2999900.000	8.09	19000.00000
2006	3000900.000	8.37	20300.00000
2007	3004200.000	8.38	20800.00000
2008	3062600.000	8.52	21800.00000
2009	3085300.000	7.83	21178.67066
2010	3090300.000	8.75	23400.00000
2011	3180600.000	9.18	23300.00000
2012	3188600.000	9.15	23800.00000
2013	3210300.000	9.62	24600.00000
2014	3225000.000	10.77	25600.00000
2015	3253000.000	11.78	26400.00000
2016	3276400.000	12.91	27200.00000
2017	3285400.000	13.24	27900.00000
2018	3305793.546	NaN	28716.39332

```
In [6]: # brief statistical description of each column
data.describe()
```

```
Out [6]:
```

	price	salesVolume	indexPrice	rate	gdp \
count	24.000000	23.000000	24.000000	24.000000	2.400000e+01

mean	256214.343276	129640.130435	63.600837	3.158854	2.305429e+06
std	126337.486812	32256.222907	31.361126	2.511183	5.698118e+05
min	74721.326880	75231.000000	18.548293	0.250000	1.336125e+06
25%	154921.017550	100414.000000	38.456497	0.500000	1.664360e+06
50%	254567.467050	132953.000000	63.192027	3.875000	2.491576e+06
75%	310237.467300	158543.500000	77.011153	5.500000	2.714552e+06
max	480240.157900	173993.000000	119.211428	7.250000	3.085300e+06

	gdpGrowth	numberHousehold	affordability	medianincome
count	24.000000	2.400000e+01	16.000000	24.000000
mean	0.021125	3.056495e+06	9.305000	20995.654520
std	0.016276	1.485513e+05	1.902546	4228.612122
min	-0.042000	2.855900e+06	6.900000	15247.902510
25%	0.017750	2.953450e+06	8.055000	17350.000000
50%	0.025000	3.002550e+06	8.635000	20550.000000
75%	0.029500	3.194025e+06	9.907500	24000.000000
max	0.043000	3.305794e+06	13.240000	28716.393320

```
In [7]: # checking which of the columns have missing values which indicated by any number > 0
data.isnull().sum()
```

```
Out[7]: price           0
salesVolume           1
indexPrice            0
rate                  0
gdp                   0
gdpGrowth             0
numberHousehold       0
affordability         8
medianincome          0
dtype: int64
```

1.2 Data Cleaning

- Fill salesVolume missing values based on a randomized values calculated from its mean and standard deviation

The inputted values will be based in the range of (mean - std) to (mean +std)

```
In [8]: salesVolume_avg = data['salesVolume'].mean()
salesVolume_std = data['salesVolume'].std()
salesVolume_null_count = data['salesVolume'].isnull().sum()
salesVolume_null_random_list = np.random.randint(salesVolume_avg - salesVolume_std, sa
data['salesVolume'][np.isnan(data['salesVolume'])] = salesVolume_null_random_list
data['salesVolume'] = data['salesVolume'].astype(int)
```

C:\ProgramData\Anaconda3\lib\site-packages\ipykernel_launcher.py:5: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html>

- Fill affordability missing values based on a randomized values calculated from its minimum value and standard deviation

The reason being is that the affordability is keep growing year by year, so we assume that the missing values which in this case is mostly the data before 2002 will be inputed by randomizing values between the range of (min - std) to (min) For the value of 2018, We used the 2017 value using the fillna function that pandas library provided using the parameter of method='ffill'

```
In [9]: data['affordability'].fillna(method='ffill',inplace=True)
        affordability_avg = data['affordability'].mean()
        affordability_std = data['affordability'].std()
        affordability_min = data['affordability'].min()
        affordability_null_count = data['affordability'].isnull().sum()
        affordability_null_random_list = np.random.randint(affordability_min - affordability_std, affordability_min, affordability_null_count)
        data['affordability'][np.isnan(data['affordability'])] = affordability_null_random_list
        data['affordability'] = data['affordability'].astype(int)
```

C:\ProgramData\Anaconda3\lib\site-packages\ipykernel_launcher.py:7: SettingWithCopyWarning: A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html>


```
In [10]: # Double checking which of the columns have missing values which indicated by any num
         data.isnull().sum()
```

```
Out[10]: price           0
         salesVolume     0
         indexPrice      0
         rate            0
         gdp             0
         gdpGrowth       0
         numberHousehold 0
         affordability    0
         medianincome     0
         dtype: int64
```


1.2.1 Data Visualisation

- Plot data to visualize the values of house price, sales volume, gdp, and number household each year.
- Visualize a correlation matrix between each feature to get a better understanding of feature importance


```
In [11]: # plot data
fig, ax = plt.subplots(figsize=(10, 5))
ax.scatter(data.price, data.index, color="b", label="House Price")
ax.set_xlabel("Value")
ax.set_ylabel("Year")
plt.show()
```




```
In [12]: fig, ax = plt.subplots(figsize=(10, 5))
ax.plot(data.price, data.index, color="r", label="salesVolume")
ax.set_xlabel("Value")
ax.set_ylabel("Year")
plt.show()
```



```
In [13]: fig, ax = plt.subplots(figsize=(10, 5))
ax.plot(data.price,data.index, color="orange", label="GDP")
ax.set_xlabel("Value")
ax.set_ylabel("Year")
plt.show()
```




```
In [14]: fig, ax = plt.subplots(figsize=(10, 5))
ax.plot(data.price,data.index, color="purple", label="numberHousehold")
ax.set_xlabel("Value")
ax.set_ylabel("Year")
plt.show()
```



In [15]: *# Build the correlation matrix*

```
matrix = data.corr()  
f, ax = plt.subplots(figsize=(10, 5))  
sns.heatmap(matrix, vmax=0.7, square=True)
```

Out[15]: <matplotlib.axes._subplots.AxesSubplot at 0x236d19ab630>



Data Preparation/Allocation

- Preparing Data by separating them into training and target (test) tables

In [16]: *# allocate all columns beside price as the training table 'train1' which will be the*
train1 = data.drop(['price'],axis=1)

In [17]: *# allocate price column as target table which will be the test*
target = data['price']

1.3 Methodology

- Gradient Boosted Regression Trees

Gradient Boosted Regression Trees (GBRT) or shorter Gradient Boosting is a flexible non-parametric statistical learning technique for classification and regression. - Linear Regression

One of the simplest model of supervised learning which assumed that the dependent variable (House Price) varies linearly with the independent variable(s) (House Sales Volume, Income, etc). Essentially fitting a straight line through the data and expecting it to give us a good prediction for values we haven't seen.

The idea of building a model is to minimize this error so that when we make a new prediction we can do so with utmost confidence (~95% is a good benchmark). There are multiple ways of minimizing this error, simplest being the least-squares method. In other words, calculating the sum of squares of each error (to eliminate negatives) and minimizing this number.

Linear Regression

```
In [18]: reg = LinearRegression()
```

```
In [19]: # split the data 90% as train 10% as test
         x_train , x_test , y_train , y_test = train_test_split(train1 , target , test_size = 0.1)
```

```
In [20]: reg.fit(x_train,y_train)
```

```
Out[20]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,
                          normalize=False)
```

```
In [21]: reg.score(x_test,y_test)
```

```
Out[21]: 0.9999999999999996
```

As we can see from the result of the LinearRegression score (~99%) which means that its way over-fitting. In order to resolve this issue, we implemented a second method, GBRT while also using GridSearchCV method to tune the parameters.

Gradient Boosting

```
In [22]: #split the data
         X_train, X_test, y_train, y_test = train_test_split(train1, target, test_size = 0.10,
```

```
In [23]: gbdt=GradientBoostingRegressor(n_estimators=10)
         gbdt.fit(X_train, y_train)
```

```
Out[23]: GradientBoostingRegressor(alpha=0.9, criterion='friedman_mse', init=None,
                                   learning_rate=0.1, loss='ls', max_depth=3, max_features=None,
                                   max_leaf_nodes=None, min_impurity_decrease=0.0,
                                   min_impurity_split=None, min_samples_leaf=1,
                                   min_samples_split=2, min_weight_fraction_leaf=0.0,
                                   n_estimators=10, n_iter_no_change=None, presort='auto',
                                   random_state=None, subsample=1.0, tol=0.0001,
                                   validation_fraction=0.1, verbose=0, warm_start=False)
```

A simple table to show the result of feature importances based on gbrt

```
In [24]: importances = gbrt.feature_importances_
std = np.std([gbrt.feature_importances_ for tree in gbrt.estimators_],
             axis=0)
indices = np.argsort(importances)[::-1]
# print the list of features and assigning number to each
count = 0
for i in list(X_train.columns.values):
    print("{} is feature {}".format(i,count))
    count +=1
print()
# Print the feature ranking
print("Feature ranking:")
for f in range(X_train.shape[1]):
    print("%d. feature %d (%f)" % (f + 1, indices[f], importances[indices[f]]))
```

```
salesVolume is feature 0
indexPrice is feature 1
rate is feature 2
gdp is feature 3
gdpGrowth is feature 4
numberHousehold is feature 5
affordability is feature 6
medianincome is feature 7
```

Feature ranking:

```
1. feature 7 (0.487810)
2. feature 1 (0.276482)
3. feature 6 (0.097341)
4. feature 5 (0.075114)
5. feature 3 (0.032181)
6. feature 2 (0.030677)
7. feature 0 (0.000395)
8. feature 4 (0.000000)
```

```
In [25]: print ("R-squared for Train: %.2f" %gbrt.score(X_train, y_train))
print ("R-squared for Test: %.2f" %gbrt.score(X_test, y_test))
```

```
R-squared for Train: 0.87
R-squared for Test: 0.79
```

Default parameters of GBRT with `n_estimators=10` already shown a good score of R-squared in both Train and Test, however, we want to see if we can optimize it more using `ShuffleSplit` cross-validation and `GridSearchCV`

- function to find the best parameter

```
In [26]: def GradientBooster(param_grid, n_jobs):

    estimator = GradientBoostingRegressor()

    #Cross-validation using ShuffleSplit which randomly shuffles and selects Train and CV
    #There are other methods like the KFold split.
    cv = ShuffleSplit(X_train.shape[0], test_size=0.1)
    #Using GridSearchCV to evaluate the classifier
    classifier = GridSearchCV(estimator=estimator, cv=cv, param_grid=param_grid, n_jobs=n_jobs)

    #We'll now fit the training dataset to this classifier
    classifier.fit(X_train, y_train)

    #print the result
    print ("Best Estimator learned through GridSearch")
    print()
    print (classifier.best_estimator_)

    return cv, classifier.best_estimator_
```

- function to plot the learning result

```
In [27]: print(__doc__)
```

```
def plot_learning_curve(estimator, title, X, y, ylim=None, cv=None,
                        n_jobs=1, train_sizes=np.linspace(.1, 1.0, 5)):
    """
    Generate a simple plot of the test and training learning curve.

    Parameters
    -----
    estimator : object type that implements the "fit" and "predict" methods
        An object of that type which is cloned for each validation.

    title : string
        Title for the chart.

    X : array-like, shape (n_samples, n_features)
        Training vector, where n_samples is the number of samples and
        n_features is the number of features.

    y : array-like, shape (n_samples) or (n_samples, n_features), optional
        Target relative to X for classification or regression;
        None for unsupervised learning.

    ylim : tuple, shape (ymin, ymax), optional
        Defines minimum and maximum yvalues plotted.
```

```

cv : integer, cross-validation generator, optional
    If an integer is passed, it is the number of folds (defaults to 3).
    Specific cross-validation objects can be passed, see
    sklearn.cross_validation module for the list of possible objects

n_jobs : integer, optional
    Number of jobs to run in parallel (default 1).
    """
plt.figure()
plt.title(title)
if ylim is not None:
    plt.ylim(*ylim)
plt.xlabel("Training examples")
plt.ylabel("Score")
train_sizes, train_scores, test_scores = learning_curve(
    estimator, X, y, cv=cv, n_jobs=n_jobs, train_sizes=train_sizes)
train_scores_mean = np.mean(train_scores, axis=1)
train_scores_std = np.std(train_scores, axis=1)
test_scores_mean = np.mean(test_scores, axis=1)
test_scores_std = np.std(test_scores, axis=1)
plt.grid()

plt.fill_between(train_sizes, train_scores_mean - train_scores_std,
                 train_scores_mean + train_scores_std, alpha=0.1,
                 color="r")
plt.fill_between(train_sizes, test_scores_mean - test_scores_std,
                 test_scores_mean + test_scores_std, alpha=0.1, color="g")
plt.plot(train_sizes, train_scores_mean, 'o-', color="r",
         label="Training score")
plt.plot(train_sizes, test_scores_mean, 'o-', color="g",
         label="Cross-validation score")

plt.legend(loc="best")
return plt

```

Automatically created module for IPython interactive environment

```

In [28]: param_grid={'n_estimators':[10],
                    'learning_rate': [0.1],# 0.05, 0.02, 0.01],
                    'max_depth':[6],#4,6],
                    'min_samples_leaf':[3],#5,9,17],
                    'max_features':[1.0],#0.3#0.1]
                }

n_jobs=4

```

```

# fit GBRT to the digits training dataset by calling the function we just created.
cv,best_est=GradientBooster(param_grid, n_jobs)

```

Best Estimator learned through GridSearch

```
GradientBoostingRegressor(alpha=0.9, criterion='friedman_mse', init=None,
                           learning_rate=0.1, loss='ls', max_depth=6, max_features=1.0,
                           max_leaf_nodes=None, min_impurity_decrease=0.0,
                           min_impurity_split=None, min_samples_leaf=3,
                           min_samples_split=2, min_weight_fraction_leaf=0.0,
                           n_estimators=10, n_iter_no_change=None, presort='auto',
                           random_state=None, subsample=1.0, tol=0.0001,
                           validation_fraction=0.1, verbose=0, warm_start=False)
```

```
In [29]: #so we got back the best estimator parameters based on GridSearch as follows:
print ("Best Estimator Parameters")
print ("-----")
print ("n_estimators: %d" %best_est.n_estimators)
print ("max_depth: %d" %best_est.max_depth)
print ("Learning Rate: %.1f" %best_est.learning_rate)
print ("min_samples_leaf: %d" %best_est.min_samples_leaf)
print ("max_features: %.1f" %best_est.max_features)

print ("Train R-squared: %.2f" %best_est.score(X_train,y_train))


#We believe that each of these parameters is critical for the method to learn better.
#We hope that some of them will help address overfitting issues as well.
```

Best Estimator Parameters

```
-----
n_estimators: 10
max_depth: 6
Learning Rate: 0.1
min_samples_leaf: 3
max_features: 1.0
Train R-squared: 0.87
```

```
In [30]: #call the plot_learning_curve module by feeding it the estimator
```

```
#The module simply runs the estimator multiple times on subsets of the data provided
#We're feeding the best parameters we've learned from GridSearchCV to the estimator
#We may need to adjust the hyperparameters further if there is overfitting or underfi
title = "Learning Curves (Gradient Boosted Regression Trees)"
estimator = GradientBoostingRegressor(n_estimators=best_est.n_estimators, max_depth=best_est.max_depth,
                                     learning_rate=best_est.learning_rate, min_samples_leaf=best_est.min_samples_leaf,
                                     max_features=best_est.max_features)
plot_learning_curve(estimator, title, X_train, y_train, cv=cv, n_jobs=n_jobs)
plt.show()
```



```
In [31]: #Re-fitting to the train set
estimator.fit(X_train, y_train)


#Calculating train/test scores - R-squared value
print ("Train R-squared: %.2f" %estimator.score(X_train, y_train))
print ("Test R-squared: %.2f" %estimator.score(X_test, y_test))
```

```
Train R-squared: 0.87
Test R-squared: 0.72
```

Looks like we've done an okay job getting about ~0.71 R-squared on the cv set. However, the default result is still better. Therefore, from the learning curve, it seems that we may be able to do a bit better with more estimators.

```
In [32]: #call the plot_learning_curve module by feeding it the estimator again but this time
title = "Learning Curves (Gradient Boosted Regression Trees)"
estimator = GradientBoostingRegressor(n_estimators=100, max_depth=best_est.max_depth,
                                     learning_rate=best_est.learning_rate, min_sample
                                     max_features=best_est.max_features)

plot_learning_curve(estimator, title, X_train, y_train, cv=cv, n_jobs=n_jobs)
plt.show()
```



```
In [33]: #Re-fitting to the train set
         estimator.fit(X_train, y_train)

         #Calculating train/test scores - R-squared value
         print ("Train R-squared: %.2f" %estimator.score(X_train, y_train))
         print ("Test R-squared: %.2f" %estimator.score(X_test, y_test))
```

```
Train R-squared: 1.00
Test R-squared: 0.97
```

It did improve the training score but there's way more overfitting. This possibly could be addressed by further reducing learning rate. However, we might not see a major improvement unless we can obtain more samples.

```
In [34]: #Switching back to the default model
         gbrt=GradientBoostingRegressor(n_estimators=10)
         gbrt.fit(X_train, y_train)
         y_pred=gbrt.predict(X_test)

         #Calculating train/test scores - R-squared value
         print ("Train R-squared: %.2f" %gbrt.score(X_train, y_train))
         print ("Test R-squared: %.2f" %gbrt.score(X_test, y_test))
```

Train R-squared: 0.87
Test R-squared: 0.79

```
In [35]: print("Prediction Result: ")
         for i in y_pred:
             print (i)

         print()
         print("Actual Result: ")
         y_test.head()
```

Prediction Result:
264021.7048683567
237654.1439182278
392462.62231642473

Actual Result:

```
Out[35]: year
         2006    251281.1776
         2005    235329.2691
         2017    480240.1579
         Name: price, dtype: float64
```

1.3.1 In conclusion, our final R-squared on the London house price dataset is ~0.79